

Artificial Intelligence

Informed Search strategies (AIMA sections 3.5 – 3.6.3 [3.5.5 and 3.5.6 excluded])

Summary

- ◇ Greedy Best-First search
- ◇ A* search
- ◇ Heuristics

Review: Tree search

```
function Tree-Search( problem, frontier ) returns a solution, or failure
  frontier ← Insert( Make-Node( problem.Initial-State ) )
  while not IsEmpty( frontier ) do
    node ← Pop( frontier )
    if problem.Goal-Test( node.State ) then return node
    frontier ← InsertAll( Expand( node, problem ) )
  end loop
  return failure
```

A strategy is defined by picking the **order of node expansion**

Best-First search

Idea: use an **evaluation function** for each node – estimate of “desirability”

⇒ Expand most desirable unexpanded node

Implementation: *frontier* is a queue sorted in decreasing order of desirability


Special cases:

- greedy best-first search

- A* search

Romania with straight-line distances to Bucharest

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



Greedy search

Evaluation function $h(n)$ (heuristic)

= estimate of cost from n to the closest goal


E.g., $h_{\text{SLD}}(n)$ = straight-line distance from n to Bucharest

Greedy search expands the node that **appears** to be closest to goal

Greedy search example

Artificial Intelligence

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374




Arad
366

Greedy search example

Artificial Intelligence

Arad 366
 Bucharest 0
 Craiova 160
 Dobreta 242
 Eforie 161
 Fagaras 176
 Giurgiu 77
 Hirsova 151
 Iasi 226
 Lugoj 244



Mehadia 241
 Neamt 234
 Oradea 380
 Pitesti 100
 Rimnicu Vilcea 193
 Sibiu 253
 Timisoara 329
 Urziceni 80
 Vaslui 199
 Zerind 374



Greedy search example

Artificial Intelligence



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



Greedy search example

Artificial Intelligence

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



Properties of greedy search

Complete??

Properties of greedy search

Complete?? No—can get stuck in loops, e.g.,

Start: Iasi, Goal: Fagaras

Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt $\rightarrow \dots$

Complete in finite space with repeated-state checking

Time??

Properties of greedy search

Complete?? No—can get stuck in loops, e.g.,

Start: Iasi, Goal: Fagaras

Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt $\rightarrow \dots$

Complete in finite space with repeated-state checking

Time?? $O(b^m)$, but a good heuristic can give dramatic improvement

Space??

Properties of greedy search

Complete?? No—can get stuck in loops, e.g.,

Start: Iasi, Goal: Fagaras

Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt $\rightarrow \dots$

Complete in finite space with repeated-state checking

Time?? $O(b^m)$, but a good heuristic can give dramatic improvement

Space?? $O(b^m)$ —keeps all nodes in memory

Optimal??

Properties of greedy search

Complete?? No—can get stuck in loops, e.g.,

Start: Iasi, Goal: Fagaras

Iasi \rightarrow Neamt \rightarrow Iasi \rightarrow Neamt $\rightarrow \dots$

Complete in finite space with repeated-state checking

Time?? $O(b^m)$, but a good heuristic can give dramatic improvement

Space?? $O(b^m)$ —keeps all nodes in memory

Optimal?? No

A* search

Idea: avoid expanding paths that are already expensive


Evaluation function $f(n) = g(n) + h(n)$

- $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost to goal from n
 - $f(n)$ = estimated total cost of path through n to goal
- ◇ A* search uses an **admissible** heuristic
i.e., $h(n) \leq h^*(n)$ where $h^*(n)$ is the **true** cost from n .
(Also require $h(n) \geq 0$, so $h(G) = 0$ for any goal G .)
- ◇ E.g., $h_{\text{SLD}}(n)$ never overestimates the actual road distance
- ◇ **Theorem:** A* search is optimal

A* search example

Artificial Intelligence

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

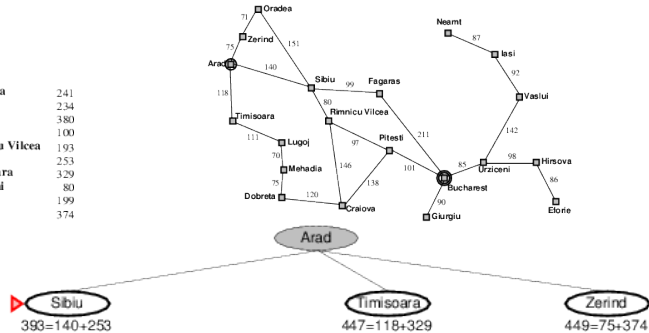


A* search example

Artificial Intelligence

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244



Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



A* search example

Artificial Intelligence



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



A* search example

Artificial Intelligence



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



A* search example

Artificial Intelligence

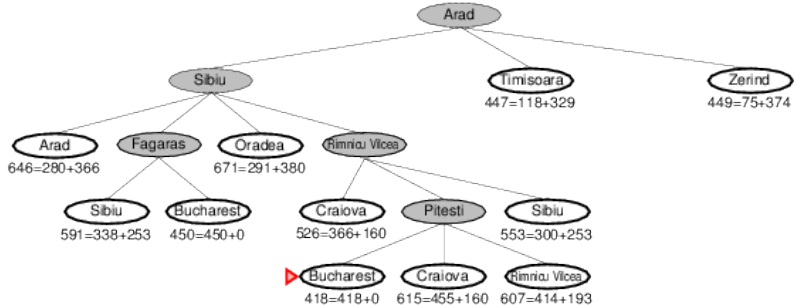

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



A* search example

Artificial Intelligence

Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Dobreta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374



Properties of A^*

Artificial
Intelligence

Complete??

Properties of A^*

Complete?? Yes, unless there are infinitely many nodes with $f \leq f(G)$
Time??

Properties of A^*

Complete?? Yes, unless there are infinitely many nodes with $f \leq f(G)$

Time?? Exponential in [relative error in $h \times$ length of soln.]

Space??

Properties of A^*

Complete?? Yes, unless there are infinitely many nodes with $f \leq f(G)$

Time?? Exponential in [relative error in $h \times$ length of soln.]

Space?? Keeps all nodes in memory

Optimal??

Properties of A^*

Complete?? Yes, unless there are infinitely many nodes with $f \leq f(G)$

Time?? Exponential in [relative error in $h \times$ length of soln.]

Space?? Keeps all nodes in memory


Optimal?? Yes, **but** it requires assumptions on heuristics (admissibility, consistency) and search strategy (tree or graph search)

Consistent heuristic definition

A heuristic is **consistent** if

$$h(n) \leq c(n, a, n') + h(n')$$

- We can show that, if h is consistent $f(n)$ is nondecreasing along any path.
- A* expands nodes in order of increasing f , hence it will find the least cost solution.



Admissible vs Consistent Heuristic

consistency \rightarrow **admissible**

Can be proved by induction on the path to goal

admissible \nrightarrow **consistency**

Find a counter example...

Tree-Search + admissible Heuristic \rightarrow optimality of A^*

Graph-Search + admissible Heuristic \nrightarrow optimality of A^*

- Can discard the optimal path to a repeated node

Graph-Search + **consistent** Heuristic \rightarrow optimality of A^*

Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$ = number of misplaced tiles

$h_2(n)$ = total **Manhattan** distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$h_1(S) = ??$

$h_2(S) = ??$

Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$ = number of misplaced tiles

$h_2(n)$ = total **Manhattan** distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(S) = ?? \quad 6$$

$$h_2(S) = ??$$

Admissible heuristics

E.g., for the 8-puzzle:

$h_1(n)$ = number of misplaced tiles

$h_2(n)$ = total **Manhattan** distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(S) = ?? \quad 6$$

$$h_2(S) = ?? \quad 4+0+3+3+1+0+2+1 = 14$$

Dominance

If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 dominates h_1 and is better for search

Typical search costs:

$d = 14$ IDS = 3,473,941 nodes

$A^*(h_1) = 539$ nodes

$A^*(h_2) = 113$ nodes

$d = 24$ IDS \approx 54,000,000,000 nodes

$A^*(h_1) = 39,135$ nodes

$A^*(h_2) = 1,641$ nodes

Given any admissible heuristics h_a, h_b ,

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible and dominates h_a, h_b

Relaxed problems

Admissible heuristics can be derived from the **exact** solution cost of a **relaxed** version of the problem

If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution

If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem

Summary

- ◇ Heuristic functions estimate costs of shortest paths
- ◇ Good heuristics can dramatically reduce search cost
- ◇ Greedy best-first search expands lowest h
 - incomplete and not always optimal
- ◇ A* search expands lowest $g + h$
 - complete and optimal
 - also optimally efficient (up to tie-breaks, for forward search)

Admissible heuristics can be derived from exact solution of relaxed problems


Exercise: Going from Lugoj to Bucharest

From Lugoj to Bucharest

- ◇ Trace the operation of A* search applied to the problem of going from Lugoj to Bucharest using the straight-line distance heuristic.
- ◇ Trace the operation of greedy best-first search applied to the problem of going from Lugoj to Bucharest using the straight-line distance heuristic.

Exercise: Navigation

Navigation with obstacles



The figure shows an artificial environment where an agent A is positioned in the square $(1, 2)^a$, the goal G is in $(3, 1)$, and there is a block B in $(2, 2)$. The agent can not pass through blocks and can move in the four directions (Up, Down, Left, Right).

^awhere the position is (row,column)

Exercise: Navigation II

Navigation with obstacles II

Formalize the problem of reaching G as a state problem

- Describe the state space, the initial and final state.
- Describe the operators.
- Find an admissible heuristics for A^* .
- Assume the operators have cost 1, draw the tree generated by A^* .

Exercise: confusing problems for greedy best-first search

confusing problems for greedy best-first search

When going from Iasi to Fagaras the straight-line distance heuristic result in poor performance for greedy best-first search. But from Fagaras to Iasi it is perfect.

Are there problems for which the heuristic produces sub-optimal paths in both directions ?